



Best practices for designing conversational agents

Autonomous vs. conversational agents	3
Best practices for conversational agents	4
Conversation design fundamentals	4
Design for conversation continuity	4
Plan for conversational breakdown	4
Establish clear conversation boundaries	4
Context and memory management	4
Design for conversation persistence	4
Handle context switches gracefully	4
Tool integration and response timing	4
Optimize for conversational pace	4
Choose tools that enhance dialogue	4
Escalation and human handoff	5
Define escalation triggers clearly	5
Design seamless transitions	5
Plan for post escalation re-engagement	5
Privacy and sensitivity management	5
Identify sensitive conversation areas	5
Design appropriate conversation boundaries	5
Evaluation and continuous improvement	5
Test diverse conversation paths	5
Monitor conversation quality indicators	5
Persona and communication style	5
Maintain consistent personality	5
Adapt communication style appropriately	6
Handle misunderstandings gracefully	6

Autonomous vs. conversational agents

The key differences between autonomous and conversational agents are as follows:

Feature	Autonomous agent	Conversational agent
Interaction model	Single-turn, task execution based on initial prompt	Multi-turn, back and forth dialogue
Use case	Automating well-defined processes triggered with a single request	Real-time user assistance, support agents, interactive information gathering
User input	Single user prompt with all inputs provided upfront	Continuous stream of chat messages
Strength	Executing complex plans with multiple tools to achieve specific goals	Maintaining context over conversation and managing interactive flows

Best practices for conversational agents

Conversation design fundamentals

Design for conversation continuity

Ask yourself, “How will users naturally progress through this conversation?” Build dialogue flows that feel intuitive and allow users to easily navigate between topics or return to previous points without losing context.

Plan for conversational breakdown

Consider “What happens when the agent doesn’t understand or when the conversation goes off track?” Design graceful recovery mechanisms that redirect users back to productive dialogue paths without frustration.

Establish clear conversation boundaries

Determine “What topics should this agent handle directly, and when should it redirect or escalate?” This prevents the agent from attempting conversations beyond its capability or authority.

Context and memory management

Design for conversation persistence

Consider “What information from this conversation should carry forward to future interactions?” Balance useful memory retention with privacy considerations and data management policies.

Handle context switches gracefully

Plan for “How should the agent respond when users change topics mid-conversation?” Build flexibility to follow natural conversation flow while maintaining helpful responses across topic boundaries.

Tool integration and response timing

Optimize for conversational pace

Ask “How long should users wait for responses, and when should the agent indicate it’s processing?” Design tool integrations that maintain natural conversation rhythm, using progress indicators for longer operations.

Choose tools that enhance dialogue

Consider “Which tools provide immediate value within conversation flow versus those that interrupt it?” Prioritize fast-responding tools like API workflows over lengthy automation processes for better conversational experience.

Escalation and human handoff

Define escalation triggers clearly

Ask “What conversation scenarios require human intervention, and how can the agent recognize these situations?” Establish clear criteria for escalating complex issues, sensitive topics, or user frustration.

Design seamless transitions

Consider “How can the agent prepare both the user and human agent for a smooth handoff?” Provide context summaries and clear expectations about what happens next in the escalation process.

Plan for post escalation re-engagement

Determine “How should the agent behave when conversations resume after human intervention?” Design continuity that acknowledges the escalation while moving forward productively.

Privacy and sensitivity management

Identify sensitive conversation areas

Ask “What topics or information require special handling, and how should the agent respond to these situations?” Build awareness of confidential information, personal data, and topics requiring human judgment.

Design appropriate conversation boundaries

Consider “How should the agent decline to discuss certain topics while maintaining helpfulness?” Create polite deflection strategies that redirect users to appropriate resources or personnel.

Evaluation and continuous improvement

Test diverse conversation paths

Ask “What different ways might users approach this conversation, including unexpected or challenging scenarios?” Build evaluation sets that cover happy paths, edge cases, and difficult interactions.

Monitor conversation quality indicators

Consider “What signals indicate successful versus unsuccessful conversations?” Track metrics like user satisfaction, task completion, and escalation rates to identify improvement opportunities.

Persona and communication style

Maintain consistent personality

Ask “How should this agent sound and behave throughout extended conversations?” Develop a clear conversational persona that remains consistent across different topics and interaction lengths.

Adapt communication style appropriately

Consider “How should the agent adjust its tone and language for different users or conversation contexts?” Balance consistency with appropriate adaptation to user needs and conversation formality.

Handle misunderstandings gracefully

Determine “How should the agent respond when users are confused or frustrated?” Design empathetic responses that acknowledge issues and guide users toward resolution.